

REINFORCEMENT LEARNING

2211cs020685

Ravuri Mourya

Alml-Zeta

```
[1]: import math
import random

class UCBGame:
    def __init__(self, num_arms):
        """
        Initialize the UCBGame instance.
        :param num_arms: Number of arms in the multi-armed bandit.
        """
        self.num_arms = num_arms
        self.counts = [0] * num_arms # Number of times each arm has been played
        self.values = [0.0] * num_arms # Average reward for each arm
        self.total_plays = 0

    def select_arm(self):
        """
        Selects the arm to play based on the Upper Confidence Bound (UCB)
        ↪algorithm.
        """
        if self.total_plays < self.num_arms:
            # Play each arm at least once
            return self.total_plays

        # Calculate UCB values for all arms
        ucb_values = [
            self.values[i] + math.sqrt((2 * math.log(self.total_plays)) / self.
            ↪counts[i])
            for i in range(self.num_arms)
        ]
        # Return the arm with the maximum UCB value
        return ucb_values.index(max(ucb_values))

    def update(self, chosen_arm, reward):
        """
        Updates the counts and values for the chosen arm based on the received
        ↪reward.
        :param chosen_arm: The arm that was chosen.
```

```

        :param reward: The reward received from the chosen arm.
        """
        self.counts[chosen_arm] += 1
        self.total_plays += 1
        n = self.counts[chosen_arm]
        value = self.values[chosen_arm]

        # Update the value using incremental formula
        self.values[chosen_arm] = ((n - 1) / n) * value + (1 / n) * reward

def simulate_game(num_arms, num_rounds, reward_probabilities):
    """
    Simulates a game where a player uses the UCB algorithm to maximize rewards.
    :param num_arms: Number of actions (arms) available to the player.
    :param num_rounds: Total number of rounds to play.
    :param reward_probabilities: List of probabilities for each arm to give a
    ↪reward.
    :return: Total rewards accumulated, counts of arm plays, and estimated
    ↪values for each arm.
    """
    game = UCBGame(num_arms)
    total_rewards = 0

    for _ in range(num_rounds):
        chosen_arm = game.select_arm()
        # Simulate reward: 1 with given probability, 0 otherwise
        reward = 1 if random.random() < reward_probabilities[chosen_arm] else 0
        game.update(chosen_arm, reward)
        total_rewards += reward

    return total_rewards, game.counts, game.values

if __name__ == "__main__":
    # Define game parameters
    num_arms = 5
    num_rounds = 1000
    reward_probabilities = [0.1, 0.2, 0.3, 0.5, 0.8] # True probabilities of
    ↪rewards

    # Run simulation
    total_rewards, counts, values = simulate_game(num_arms, num_rounds,
    ↪reward_probabilities)

    # Output results
    print(f"Total rewards: {total_rewards}")

```

```
print(f"Number of times each arm was played: {counts}")
print(f"Estimated values for each arm: {values}")
```

Total rewards: 703

Number of times each arm was played: [19, 27, 55, 76, 823]

Estimated values for each arm: [0.05263157894736842, 0.18518518518518517, 0.39999999999999998, 0.47368421052631576, 0.7764277035236941]

```
[2]: import math
import random

class UCBOptimizer:
    def __init__(self, num_modes):
        """
        Initializes the UCBOptimizer.
        :param num_modes: Number of possible settings (modes) for the devices.
        """
        self.num_modes = num_modes # Number of possible settings for devices
        self.counts = [0] * num_modes # Number of times each mode has been
        ↪selected
        self.values = [0.0] * num_modes # Average efficiency for each mode
        self.total_selections = 0 # Total number of selections made

    def select_mode(self):
        """
        Selects the mode to use based on the Upper Confidence Bound (UCB)
        ↪algorithm.
        """
        if self.total_selections < self.num_modes:
            # Ensure each mode is selected at least once
            return self.total_selections

        # Calculate UCB values for each mode
        ucb_values = [
            self.values[i] + math.sqrt((2 * math.log(self.total_selections)) /
        ↪self.counts[i])
            for i in range(self.num_modes)
        ]

        # Return the mode with the highest UCB value
        return ucb_values.index(max(ucb_values))

    def update(self, chosen_mode, efficiency):
        """
        Updates the counts and values for the chosen mode based on the observed
        ↪efficiency.
        :param chosen_mode: The mode that was chosen.
```

```

        :param efficiency: The observed efficiency value.
        """
        self.counts[chosen_mode] += 1
        self.total_selections += 1
        n = self.counts[chosen_mode]
        value = self.values[chosen_mode]

        # Update the average efficiency using the incremental formula
        self.values[chosen_mode] = ((n - 1) / n) * value + (1 / n) * efficiency

def simulate_smart_home(num_modes, num_rounds, efficiency_factors):
    """
    Simulates an IoT smart home system optimizing energy usage across multiple
    ↪ modes.
    :param num_modes: Number of settings (modes) available for devices.
    :param num_rounds: Total number of iterations for optimization.
    :param efficiency_factors: List of real efficiency probabilities for each
    ↪ mode.
    :return: Total efficiency score and mode selection statistics.
    """
    optimizer = UCB0Optimizer(num_modes)
    total_efficiency = 0

    for _ in range(num_rounds):
        chosen_mode = optimizer.select_mode()

        # Simulate efficiency: value based on real efficiency factor plus
        ↪ random noise
        efficiency = efficiency_factors[chosen_mode] + random.uniform(-0.05, 0.
        ↪ 05)
        efficiency = max(0, min(1, efficiency)) # Clamp efficiency between 0
        ↪ and 1

        optimizer.update(chosen_mode, efficiency)
        total_efficiency += efficiency

    return total_efficiency, optimizer.counts, optimizer.values

if __name__ == "__main__":
    # Define simulation parameters
    num_modes = 4 # Example: Low, Medium, High, Auto settings for devices
    num_rounds = 500 # Number of rounds to simulate
    efficiency_factors = [0.7, 0.8, 0.9, 0.85] # Base efficiency for each mode

    # Run the simulation

```

```

    total_efficiency, counts, values = simulate_smart_home(num_modes,
↳ num_rounds, efficiency_factors)

    # Display results
    print(f"Total efficiency score: {total_efficiency:.2f}")
    print(f"Number of times each mode was selected: {counts}")
    print(f"Estimated efficiencies for each mode: {values}")

```

Total efficiency score: 421.84

Number of times each mode was selected: [62, 101, 202, 135]

Estimated efficiencies for each mode: [0.7020655744272584, 0.8017331823941076, 0.9048560301754909, 0.8485900024313898]

```

[3]: import random
import math

class PACChessGame:
    def __init__(self, board_size, max_depth, epsilon, delta):
        """
        Initialize the PAC Chess game.
        :param board_size: Size of the chessboard.
        :param max_depth: Number of moves ahead to evaluate.
        :param epsilon: Error tolerance for PAC decision-making.
        :param delta: Confidence threshold for PAC decision-making.
        """
        self.board_size = board_size
        self.max_depth = max_depth
        self.epsilon = epsilon
        self.delta = delta
        self.board = self.initialize_board()

    def initialize_board(self):
        """Initializes a simplified chessboard."""
        board = [[None for _ in range(self.board_size)] for _ in range(self.
↳ board_size)]

        # Place pawns for both players
        for i in range(self.board_size):
            board[1][i] = 'P1' # Player 1 pawns
            board[-2][i] = 'P2' # Player 2 pawns

        # Place rooks
        board[0][0], board[0][-1] = 'R1', 'R1' # Player 1 rooks
        board[-1][0], board[-1][-1] = 'R2', 'R2' # Player 2 rooks

        # Place kings
        board[0][self.board_size // 2] = 'K1' # Player 1 king

```

```

board[-1][self.board_size // 2] = 'K2' # Player 2 king

return board

def possible_moves(self, player):
    """Generates all possible moves for the given player."""
    moves = []
    for row in range(self.board_size):
        for col in range(self.board_size):
            if self.board[row][col] and self.board[row][col].
↪endswith(player):
                # Example: Simplified pawn movement
                if row + 1 < self.board_size:
                    moves.append((row, col), (row + 1, col)) # Move_
↪forward
    return moves

def evaluate_board(self):
    """Evaluates the board state for scoring."""
    score = 0
    for row in self.board:
        for piece in row:
            if piece == 'P1':
                score += 1
            elif piece == 'P2':
                score -= 1
            elif piece == 'K1':
                score += 10
            elif piece == 'K2':
                score -= 10
    return score

def pac_decision(self, moves):
    """Selects the best move using PAC principles."""
    best_move = None
    best_score = float('-inf')

    # Calculate the number of samples needed for PAC guarantee
    samples = int((1 / (self.epsilon ** 2)) * math.log(1 / self.delta))

    for move in moves:
        total_score = 0
        for _ in range(samples):
            total_score += self.simulate_move(move)
        avg_score = total_score / samples

        if avg_score > best_score:

```

```

        best_score = avg_score
        best_move = move

    return best_move

def simulate_move(self, move):
    """Simulates the effect of a move and evaluates the board."""
    start, end = move
    piece = self.board[start[0]][start[1]]

    # Make the move
    self.board[start[0]][start[1]] = None
    self.board[end[0]][end[1]] = piece

    # Evaluate the board
    score = self.evaluate_board()

    # Undo the move
    self.board[end[0]][end[1]] = None
    self.board[start[0]][start[1]] = piece

    return score

def play_game(self):
    """Plays a game between two PAC-based players."""
    current_player = '1'

    for turn in range(50): # Maximum number of turns
        moves = self.possible_moves(current_player)
        if not moves:
            print(f"Player {current_player} has no moves left. Game over.")
            break

        chosen_move = self.pac_decision(moves)
        print(f"Player {current_player} chooses move {chosen_move}")

        start, end = chosen_move
        piece = self.board[start[0]][start[1]]
        self.board[start[0]][start[1]] = None
        self.board[end[0]][end[1]] = piece

        current_player = '2' if current_player == '1' else '1'

    print("Game ended.")

if __name__ == "__main__":

```

```
# Define game parameters
```

```
game = PACChessGame(board_size=6, max_depth=3, epsilon=0.1, delta=0.05)  
game.play_game()
```

```
Player 1 chooses move ((0, 0), (1, 0))  
Player 2 chooses move ((4, 3), (5, 3))  
Player 1 chooses move ((0, 3), (1, 3))  
Player 2 chooses move ((4, 0), (5, 0))  
Player 1 chooses move ((0, 5), (1, 5))  
Player 2 chooses move ((4, 1), (5, 1))  
Player 1 chooses move ((1, 0), (2, 0))  
Player 2 chooses move ((4, 2), (5, 2))  
Player 1 chooses move ((1, 1), (2, 1))  
Player 2 chooses move ((4, 4), (5, 4))  
Player 1 chooses move ((1, 2), (2, 2))  
Player 2 chooses move ((4, 5), (5, 5))  
Player 1 chooses move ((1, 3), (2, 3))  
Player 2 has no moves left. Game over.  
Game ended.
```

```
[ ]:
```